

Download the jupyter notebook from ILIAS

or

run the notebook in colab:

<https://colab.research.google.com/drive/1eWA7jl4R7JMZFLzRMAXI5ckCKZyBKSRQ>

(<https://tinyurl.com/yx2xnook>)



ATML Tutorial 03

Advanced Topics in Machine Learning

03.03.2020

Adam Bielski



Information

- First assignment now on ILIAS (due date: 24/03/2020)
- Bring your laptops next week



Last week

- Basic PyTorch (**tensors**, **CPU/GPU** computations, **gradients** with torch.autograd)
- Feeding data with PyTorch **Dataset** objects (`__getitem__`, `__len__`)
- Iterating datasets with **DataLoaders**
- Creating **models with nn.Module** subclasses (forward method) and linear layers (nn.Linear)
- Running the models



Today

- training neural networks in PyTorch
- debugging and analyzing overfitting/underfitting



Neural networks

- Input (features): \mathbf{x} *pixels of cat/dogs images*
- Output (labels/target values): \mathbf{y} *cat/dog label*
- We want to find a function \mathbf{f} that will approximate $\mathbf{f}(\mathbf{x}) = \mathbf{y}$ and can be run on new data when \mathbf{y} is unknown

Neural networks approximate the function \mathbf{f} by composition of functions

$$f(x) = f^{(3)}(f^{(2)}(f^{(1)}(x)))$$



Multilayer Perceptron (MLP)

$$\text{MLP: } f(x) = f^{(3)}(f^{(2)}(f^{(1)}(x)))$$


A class of neural networks, where $f^{(i)}(x) = \phi(x\mathbf{w}^T + b)$

\mathbf{w} , b - parameters we want to optimize

ϕ - non-linear function



Multilayer Perceptron (MLP)





Loss functions - how different is prediction from target values?

We will **optimize** parameters of a neural network to **minimize** a loss function.

Problem	Output activation	Loss function
Regression y = real value	Linear / no activation	Mean square error $L = \text{mean}([y - y']^2)$
Regression (within range) y scaled to y = [-1; 1]	Tanh [-1; 1] / Sigmoid [0; 1]	Mean square error $L = \text{mean}([y - y']^2)$
Binary classification y = 0 or y = 1	Sigmoid $1/(1+\exp(-x))$ Probability of class y=1	Binary cross-entropy $L = -y \cdot \log(y') - (1-y) \cdot \log(1-y')$
Multi-class classification y=0..n	Softmax $y' = \exp(x_i) / \sum(\exp_j)$ Vector of probabilities for each class	Cross-entropy $L = -\log(y'[y])$ (maximize probability of true class)



Loss functions - PyTorch

PyTorch documentation - <https://pytorch.org/docs/stable/nn.html#loss-functions>

Problem	Output activation (in model)	Loss function
Regression y = real valued vector	<code>nn.Linear(h_dim, y_dim)</code>	<code>nn.MSELoss()</code>
Regression (within range) y scaled to y = [-1; 1]	<code>nn.Linear(h_dim, y_dim),</code> <code>nn.Tanh()</code>	<code>nn.MSELoss()</code>
Binary classification y = 0 or y = 1	<code>nn.Linear(h_dim, 1),</code> <code>nn.Sigmoid()</code>	<code>nn.BCELoss()</code>
	<code>nn.Linear(h_dim, 1)</code>	<code>nn.BCEWithLogitsLoss()</code> // combines Sigmoid and NLLLoss
Multi-class classification y=0..n	<code>nn.Linear(h_dim, n),</code> <code>nn.LogSoftmax()</code>	<code>nn.NLLLoss()</code>
	<code>nn.Linear(h_dim, n)</code>	<code>nn.CrossEntropyLoss()</code> // combines LogSoftmax and NLLLoss



Training neural networks (high-level description)

Stochastic Gradient Descent

Input: feature vectors $x=\{x_i\}$, targets $y=\{y_i\}$,
neural network model $f(x; \theta)$ with randomly initialized parameters θ

Output: optimized model parameters θ

for number of training iterations **do**

sample minibatch of m samples $\{(x_i, y_i)\}$

 compute **predicted** $y'=f(x; \theta)$

 compute **loss** $L=loss_function(y', y)$

 compute **gradients** $\nabla_{\theta} L$


 update parameters $\theta \leftarrow \theta - \alpha \nabla_{\theta} L$

// backpropagation algorithm

// update parameters to minimize loss

Training neural networks (high-level description)

Stochastic Gradient Descent



Source: <https://nikcheerla.github.io/deeplearningschool/2017/09/08/ch2.1-linear-regression-sgd/>

Let's train some networks!

Cat VS Dog



Like



Love