



# ATML Tutorial 05


## Regularization

Advanced Topics in Machine Learning  
17.03.2020  
Adam Bielski

# Overfitting


## Tutorial 3: dogs vs cats classification

Linear model



**Underfitting: model is too simple to get good results**

MLP with 2 hidden layers



**Overfitting: Model is too complex for our dataset**

- memorizes training examples
- doesn't generalize to new data



# Overfitting

- We want to use complex models that can learn better representations
- We can use regularization to reduce overfitting





# Regularization

Option A: get more data => expensive, difficult  
generate more data => **data augmentation**

Option B: constrain your model  
- make it more difficult to memorize the data  
=> **L2 regularization, Dropout, Early stopping**

# Data augmentation


Apply transformations to the data that preserves a label,  
e.g. Crop, Rotation, Affine Transformation, Color Jitter...



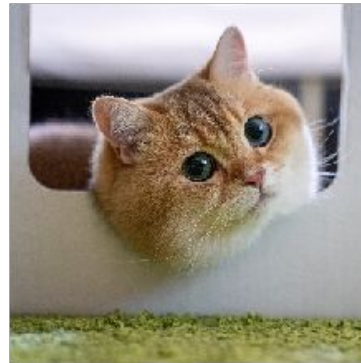

STILL A CAT

# Data augmentation

Choose carefully, data-domain specific!




Horizontal Flip



**STILL A CAT**

- Label preserved



**NOT A FIVE ANYMORE**

- Label changed



# Data augmentation - PyTorch

Many of them implemented in **torchvision** package

<https://pytorch.org/docs/stable/torchvision/transforms.html>

See: `augmentation_transforms.ipynb`

Google Colab:

<https://colab.research.google.com/drive/1VUMkwTyubaaFC4lyL6FzOikJRSGu36Tt>

<https://tinyurl.com/uy3ptv7>



# Data augmentation - PyTorch

Example with ImageFolder dataset class

```
CLASS torchvision.datasets.ImageFolder(root, transform=None, target_transform=None, loader=  
<function default_loader>)
```


```
from torchvision.datasets import ImageFolder  
from torchvision.transforms import Resize, ToTensor, Normalize,  
Compose  
  
target_size = (32, 32)  
transforms = Compose([Resize(target_size), # Resizes image  
                     ToTensor(),          # Converts to Tensor  
                     Normalize(mean=(0.5, 0.5, 0.5),  
                               std=(0.5, 0.5, 0.5)), # scaling  
                     ])  
  
train_dataset = ImageFolder(data_dir, transform=transforms)
```

```
from torchvision.datasets import ImageFolder  
from torchvision.transforms import Resize, ToTensor, Normalize,  
Compose  
  
transforms = Compose([Resize((40, 40)), # Resizes image  
                     RandomCrop((32, 32)), # Crop 32x32 area  
                     RandomHorizontalFlip(),  
                     ToTensor(),          # Converts to Tensor  
                     Normalize(mean=(0.5, 0.5, 0.5),  
                               std=(0.5, 0.5, 0.5)),  
                     ])  
  
train_dataset = ImageFolder(data_dir, transform=transforms)
```



# Early stopping

Monitor validation loss / accuracy, save the model with the best value




● is better than ● ;

Monitor loss and get the model from ●

Stop the training if the validation loss starts increasing

# Early stopping

Monitor validation loss / accuracy, save the model with the best value



```
best_val_loss = np.inf
best_model = None
patience = 5 # if no improvement after 5 epochs, stop tra
counter = 0
```

```
for epoch in range(n_epochs):
    ### Train for an epoch and evaluate on validation set
    ### (...)

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_model = deepcopy(model)
        counter = 0
    else:
        counter += 1
    if counter == patience:
        print('No improvement for {} epochs;
              training stopped.'.format(patience))
        break
```



## L2 Regularization / weight decay

Constrain magnitude of **weights** / **parameters** of the neural network by adding a term to the loss function

$$L = L_{cross-entropy} + \boxed{\frac{1}{2} \alpha \sum_i \theta_i^2}$$

L2 regularization

Theta - neural network weights

Alpha - regularization strength (if too big, all weights will shrink to 0 - network will not learn)

In PyTorch we can define it in the optimizer:

```
CLASS torch.optim.SGD(params, lr=<required parameter>, momentum=0, dampening=0,  
    weight_decay=0, nesterov=False)
```


[SOURCE]

^ We provide *alpha* here


```
alpha = 0.001  
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate, weight_decay=alpha)
```

# Dropout

- Set hidden layer outputs to zero in each iteration randomly; for training
- Use all the outputs (do nothing) for evaluation



(a) Standard Neural Network



(b) Network after Dropout



# Dropout

In PyTorch - we add `nn.Dropout()` layers in model definition

```
class MLPModelDropout(nn.Module):  
    // probability of setting a unit to 0  
  
    def __init__(self, input_dim, hidden_dim, dropout_p=0.5):  
        super(MLPModelDropout, self).__init__()  
        self.layers = nn.Sequential(  
            nn.Linear(input_dim, hidden_dim),  
            nn.ReLU(),  
            nn.Dropout(dropout_p), ### Adding dropout layer  
            nn.Linear(hidden_dim, hidden_dim),  
            nn.ReLU(),  
            nn.Dropout(dropout_p), ### Adding dropout layer  
            nn.Linear(hidden_dim, hidden_dim),  
            nn.ReLU(),  
            nn.Dropout(dropout_p), ### Adding dropout layer  
            nn.Linear(hidden_dim, 2)  
        )  
  
    def forward(self, input):  
        input = input.view(input.size(0), -1)  
        return self.layers(input)
```



# Dropout

```
model = MLPModelDropout(32*32*3, 128, 0.5)
model.train() # Sets all the layers to training mode - dropout is active
model.eval()  # Seta all the layers to evaluation mode - dropout is inactive (does nothing)
```

**See `dropout_layer.ipynb`**